

MM916 Additional tidyverse questions

The following questions are based on data from the RStats community activity TidyTuesday. A new data set is released each week and the goal is to tease out interesting facts about the data. Some pretty spectacular data visualisations have been part of this... search for #tidytuesday on twitter/X.

Note that the instructions here are deliberately vague. You'll have to work out how to find out each piece of information: this is good practice for doing real-life analysis!

Question 1

- a) Load the data stored in `volcano.RData`.

```
# use the load() function to read in RData files
load("volcano.RData")
```

- b) Which country has had the most recent volcanic eruption?

```
# There are a couple of options to do this
```

```
# Option 1: Use volcanos - need to extract "unknown"
# Can get country directly
volcano %>%
  filter(last_eruption_year!="Unknown") %>%
  mutate(last_eruption_year=as.numeric(last_eruption_year)) %>%
  arrange(desc(last_eruption_year))
```

```
## # A tibble: 657 x 26
##   volcano_number volcano_name   primary_volcano_type last_eruption_year country
##           <dbl> <chr>          <chr>                      <dbl> <chr>
## 1         282080 Aira            Caldera                   2020 Japan
## 2         282110 Asosan          Caldera                   2020 Japan
## 3         255020 Bagana          Lava cone                2020 Papua ~
## 4         300250 Bezymianny     Stratovolcano             2020 Russia
## 5         357070 Chillan, Neva~ Stratovolcano             2020 Chile
## 6         268010 Dukono          Complex                  2020 Indone~
## 7         290380 Ebeko           Stratovolcano             2020 Russia
## 8         390020 Erebus          Stratovolcano             2020 Antarc~
## 9         221080 Erta Ale        Shield                   2020 Ethiop~
## 10        211060 Etna            Stratovolcano(es)       2020 Italy
## # i 647 more rows
## # i 21 more variables: region <chr>, subregion <chr>, latitude <dbl>,
## #   longitude <dbl>, elevation <dbl>, tectonic_settings <chr>,
## #   evidence_category <chr>, major_rock_1 <chr>, major_rock_2 <chr>,
## #   major_rock_3 <chr>, major_rock_4 <chr>, major_rock_5 <chr>,
## #   minor_rock_1 <chr>, minor_rock_2 <chr>, minor_rock_3 <chr>,
## #   minor_rock_4 <chr>, minor_rock_5 <chr>, population_within_5_km <dbl>, ...
```

```
# Option 2: (Better - more detail) Use eruptions and order by the three start_ columns
# Turns out it already is!
```

```

# need to join with volcanos to get country
eruptions %>%
  arrange(desc(start_year), desc(start_month), desc(start_day)) %>%
  right_join(volcano) %>%
  # can select columns to make lookup easier include date for sanity check
  select(volcano_number, volcano_name, country, start_year:start_day)

## # A tibble: 9,828 x 6
##   volcano_number volcano_name      country start_year start_month start_day
##   <dbl> <chr>          <chr>        <dbl>       <dbl>       <dbl>
## 1 266030 Soputan     Indone~        2020         3       23
## 2 233020 Fournaise, Piton de ~ France        2020         2       10
## 3 345020 Rincon de la Vieja Costa ~        2020         1       31
## 4 273070 Taal        Philip~        2020         1       12
## 5 282050 Kuchinoerabujima Japan        2020         1       11
## 6 241040 Whakaari/White Island New Ze~    2019        12       9
## 7 311060 Semisopochnoi United~        2019        12       7
## 8 282060 Kikai        Japan        2019        11       2
## 9 300260 Klyuchevskoy Russia       2019        10      24
## 10 283110 Asamayama   Japan        2019         8       7
## # i 9,818 more rows

# Option 3: (Even better - which has been erupting most recently) Use eruptions
# and order by end_ columns
# need to join with volcanos to get country
eruptions %>%
  arrange(desc(end_year), desc(end_month), desc(end_day)) %>%
  right_join(volcano) %>%
  # can select columns to make lookup easier include date for sanity check
  select(volcano_number, volcano_name, country, end_year:end_day)

## # A tibble: 9,828 x 6
##   volcano_number volcano_name      country   end_year end_month end_day
##   <dbl> <chr>          <chr>        <dbl>       <dbl>       <dbl>
## 1 390020 Erebus        Antarctica  2020         4       18
## 2 345020 Rincon de la Vieja Costa Rica  2020         4       17
## 3 282050 Kuchinoerabujima Japan        2020         4       17
## 4 300260 Klyuchevskoy Russia       2020         4       17
## 5 282110 Asosan        Japan        2020         4       17
## 6 352090 Sangay         Ecuador      2020         4       17
## 7 262000 Krakatau     Indonesia    2020         4       17
## 8 263250 Merapi        Indonesia    2020         4       17
## 9 261170 Kerinci      Indonesia    2020         4       17
## 10 282080 Aira          Japan        2020         4       17
## # i 9,818 more rows

```

c) How many eruptions are recorded as having a severity index greater than 4?

```

# You will need to filter out the missing values
eruptions %>% filter(!is.na(vei)) %>%
  summarise(severe_eruptions=sum(vei>4))

```

```

## # A tibble: 1 x 1
##   severe_eruptions
##   <int>
## 1 237

```

```
#Note that another way to do the same thing is  
eruptions %>% summarise(severe_eruptions=sum(vei>4,na.rm=TRUE))
```

```
## # A tibble: 1 x 1  
##   severe_eruptions  
##       <int>  
## 1             237
```

- d) Which type of volcano has produced the most severe eruption in the data set?

```
# If you interpret this literally, as the most severe single eruption, then:
```

```
eruptions %>% filter(vei==max(vei,na.rm=TRUE)) %>%
```

```
# this gives several different eruptions which are all tied for maximum VEI.
```

```
# then look up the type
```

```
left_join(volcano) %>% select(primary_volcano_type,vei)
```

```
## # A tibble: 8 x 2  
##   primary_volcano_type    vei  
##   <chr>                  <dbl>  
## 1 Stratovolcano          7  
## 2 Stratovolcano          7  
## 3 Stratovolcano          7  
## 4 Shield(s)              7  
## 5 Caldera                7  
## 6 Caldera                7  
## 7 Caldera                7  
## 8 Caldera                7
```

```
# Alternatively, if you interpret it as the volcano type with the most severe  
# eruptions *on average*:
```

```
# associate type with each individual eruption
```

```
eruptions %>% left_join(volcano) %>%
```

```
# find the mean severity for each type (as in 3c, have to eliminate NA values)
```

```
group_by(primary_volcano_type) %>% summarise(mean_vei = mean(vei,na.rm=TRUE)) %>%
```

```
# finally, sort by mean VEI so that the most severe pop to the top of the data frame
```

```
arrange(desc(mean_vei))
```

```
## # A tibble: 24 x 2  
##   primary_volcano_type  mean_vei  
##   <chr>                  <dbl>  
## 1 Caldera(s)            3.38  
## 2 Subglacial             2.82  
## 3 Maar(s)                2.67  
## 4 Pyroclastic cone      2.6  
## 5 Compound                2.57  
## 6 Pyroclastic shield     2.46  
## 7 Shield(s)               2.42  
## 8 Stratovolcano?         2.38  
## 9 Lava dome(s)            2.35  
## 10 Fissure vent(s)        2.27  
## # i 14 more rows
```

- e) On average, does there appear to be an effect of major rock type 1 on the severity of an eruption?

```
# Need to join the two data sets
```

```
# use left join since eruptions is bigger
```

```
eruptions %>%
```

```

left_join(volcano) %>%
# Want to find the average severity based on major rock type 1
group_by(major_rock_1) %>%
# I've used median as the average since it is more robust
summarise(median(vei, na.rm=TRUE))

## # A tibble: 11 x 2
##   major_rock_1      `median(vei, na.rm = TRUE)`<dbl>
##   <chr>                      2
## 1 Andesite / Basaltic Andesite    2
## 2 Basalt / Picro-Basalt          2
## 3 Dacite                         2
## 4 Foidite                         1
## 5 Phono-tephrite / Tephri-phonolite 3
## 6 Phonolite                       2
## 7 Rhyolite                        2
## 8 Trachyandesite / Basaltic Trachyandesite 2
## 9 Trachybasalt / Tephrite Basanite    2
## 10 Trachyte / Trachydacite        3
## 11 <NA>                           2

# certain rock types do appear to have different average severity indices.
# Not enough information here to say with any certainty but trachyte and
# phono-tephrite appear to have slightly higher severity and foidite has lower.
# Majority seem to have a median severity index of 2.

```

Question 2

- a) Load the data stored in plants1.RData

```
load("plants1.RData")
```

- b) Find out how many species are extinct versus extinct in the wild.

```

plants %>%
  group_by(red_list_category) %>%
  # The function n() takes a count
  summarise(count=n())

```

```

## # A tibble: 2 x 2
##   red_list_category   count
##   <chr>                 <int>
## 1 Extinct                435
## 2 Extinct in the Wild     65

```

- c) Of the species that are threatened, how many have had action taken in 2 or fewer areas?

```

# Option 1: Simply add up the columns
plants %>%
  # Get the sum of the action columns
  mutate(ActionsTaken = action_LWP + action_SM + action_LP + action_RM + action_EA + action_NA) %>%
  summarise(sum(ActionsTaken<=2))

## # A tibble: 1 x 1
##   `sum(ActionsTaken <= 2)`<int>
##   <int>
## 1 65

```

```

## 1 485
# Option 2: Not covered in class but an additional option is c_across()
plants %>%
  rowwise() %>%
  # Get the sum of the action columns
  mutate(ActionsTaken=sum(c_across(action_LWP:action_NA))) %>%
  ungroup() %>%
  summarise(sum(ActionsTaken<=2))

## # A tibble: 1 x 1
##   `sum(ActionsTaken <= 2)`
##   <int>
## 1 485
# In option 2 rowwise() indicates that we want to do operations row-wise rather
# than column-wise - by default sum() would want to take the sum of the columns.
# then c_across allows us to specify column names in the same way as normal.
# rowwise() is a special type of grouping - to go back to standard use use ungroup()

```

- d) Which country contains the most threatened or extinct species of plant?

```

plants %>%
  group_by(country) %>%
  summarise(Count=n()) %>%
  arrange(desc(Count))

## # A tibble: 72 x 2
##   country     Count
##   <chr>      <int>
## 1 Madagascar  98
## 2 United States 66
## 3 Ecuador    52
## 4 Tanzania   25
## 5 Malaysia   18
## 6 Burundi    17
## 7 Guinea     14
## 8 Indonesia   12
## 9 New Caledonia 11
## 10 South Africa 11
## # i 62 more rows
# Madagascar has the most followed by the US and Ecuador

```

- e) Which plant type is the most threatened?

```

# if "threatened" means "on this list but not yet fully extinct"
plants %>% filter(red_list_category != "Extinct") %>%
  group_by(group) %>%
  summarise(Count=n()) %>%
  arrange(desc(Count))

## # A tibble: 3 x 2
##   group     Count
##   <chr>      <int>
## 1 Flowering Plant  60
## 2 Cycad        4
## 3 Ferns and Allies 1

```

```

# flowering plants are by far the most commonly observed in the data.
# The next question, however, is: What is the proportion of known flowering plant
# species relative to all plants? If there are more flowering plants than any other
# kind then this result might be expected...if not then that suggests the flowering
# plants are more easily threatened than other types of plant.`

```

- f) During which time period did most plants go extinct? (Exclude plants that went extinct pre-1900 and those with no registered date.)

```

plants %>%
  # exclude data that we don't need
  filter(!year_last_seen %in% c("Before 1900", NA)) %>%
  group_by(year_last_seen) %>%
  summarise(Count=n()) %>%
  arrange(desc(Count))

```

```

## # A tibble: 6 x 2
##   year_last_seen Count
##   <chr>           <int>
## 1 1940-1959        74
## 2 1900-1919        70
## 3 1920-1939        70
## 4 1960-1979        60
## 5 2000-2020        52
## 6 1980-1999        44

#1940-1959 had the most extinctions. The trend seems to be largely decreasing
# up until the most recent 20 year period.`

```

Question 3

Read in the data stored in `rap_rankings.RData`

The algorithm used to rank the rap artists is described at <https://github.com/rfordatascience/tidytuesday/blob/master/data/2020/2020-04-14/readme.md>

Have a go at implementing the algorithm described.

```

load("rap_rankings.RData")

# this is a hard one! The sample solution below uses tidyverse functions exclusively,
# but another approach would be to go through the rows with a for loop and use if-else
# statements to keep a tally of rankings and points.

polls %>%
  # Group by all of the variables that appear in the final data set
  group_by(title, artist, gender, year, rank) %>%
  summarise(n=n()) %>%
  # we're done with the groups so we should ungroup
  ungroup() %>%
  # now have a count of how many times each song got each rank.
  # but at this point, each row = a song getting a certain ranking;
  # next, pivot_wider so that each row = a song.
  # specify values_fill to fill in zeros
  pivot_wider(names_from="rank", values_from="n", names_prefix="n", values_fill=0) %>%

```

```

# calculate the number of points and number of ranks given
mutate(points=10*n1+8*n2+6*n3+4*n4+2*n5,
      n=n1+n2+n3+n4+n5) %>%
# reorder columns
select(title, gender, points, n, n1, n2, n3, n4, n5) %>%
# Order rows: points is most important, followed by total rankings then each ranking
arrange(desc(points), desc(n),
        desc(n1), desc(n2),
        desc(n3), desc(n4), desc(n5)) %>%
rowid_to_column(var="ID")

## # A tibble: 311 x 11
##   ID title      artist gender points     n    n1    n2    n3    n4    n5
##   <int> <chr>     <chr> <chr> <dbl> <int> <int> <int> <int> <int> <int>
## 1 1 Juicy      The N~ male    140    18     9     3     3     1     2
## 2 2 Fight The Pow~ Publi~ male    100    11     7     3     1     0     0
## 3 3 Shook Ones (P~ Mobb ~ male    94     13     4     5     1     1     2
## 4 4 The Message Grand~ male    90     14     5     3     1     0     5
## 5 5 Nuthin' But A~ Dr Dr~ male    84     14     2     4     2     4     2
## 6 6 C.R.E.A.M. Wu-Ta~ male    62     10     3     1     1     4     1
## 7 7 93 'Til Infin~ Souls~ male    50     7     2     2     2     0     1
## 8 8 Passin' Me By The P~ male    48     6     3     2     0     0     1
## 9 9 N.Y. State Of~ Nas     male    46     7     1     3     1     1     1
## 10 10 Dear Mama 2Pac     male    42     6     2     1     1     2     0
## # i 301 more rows

```